

Lecture 5

Monday, October 7, 2024 10:01 AM

Housekeeping

Remind me to turn on zoom, share screen, start recording!

- HW 1 feedback coming soon
- HW 2 posted; due Friday 10/11 (20% penalty for Saturday night)
 - Remember to mark which problem is which on pdf
 - Remember to also hand in text file with executable SQL queries
- Project Part 1 due Friday 10/18
 - Teams should notify TAs by 10/11
- HW 3 coming soon
- Next week:
 - Mon 10/14 is Fall Break/Indigenous People's Day
 - Tues 10/15 is NYU "Legislative Monday"
 - I will post recorded lecture instead of in-class lecture
 - Office hours schedule TBD

Today:

More SQL

- Finish Chapter 3
- Chapter 4

Where are we?

Monday, February 27, 2023 10:04 AM

SELECT <attribute-list possibly including functions and/or aggregations>
FROM <table | join of tables | cartesian product of tables >
WHERE <predicate on individual rows (possibly involving ...IN subquery)>
GROUP BY <attribute-list including all un-aggregated attributes from select clause>
HAVING < predicate involving results of aggregation>
ORDER BY

NATURAL



Recap: Execution of SQL Queries

```
select dept_name, avg(salary)
from instructor join department
where budget > 100000
group by dept_name
having avg (salary) > 42000;
```

Is executed in the following steps:

- **from:** do the cross product or join (in his case) of the tables in the **from** clause.
- **where:** apply the condition in the **where** clause to every tuple of the table produced by the **from** clause.
- **group by:** aggregate and reorg table into a new table containing only the **group by** attributes and any aggregates referenced in the **having** or **select** clauses.
- **having:** apply condition in the **having** clause to every tuple of the table produced by the **group by**.
- **select:** project onto the requested attributes

Set Theory operations

Tuesday, March 10, 2020 9:38 PM

Set theory operation

Union

\cup

Intersection

\cap

Set Difference

$-$

Belongs to

\in

Does not belong to

\notin

Is empty

$= \emptyset$

Is not empty

$\neq \emptyset$

Comparison to each element

SQL Predicate

UNION [can often be done with pred1 OR pred2 in the WHERE clause, instead]

INTERSECT [Not supported in MySQL]

EXCEPT [MINUS in Oracle; not supported in MySQL]

IN (subquery describing X)

\subseteq

NOT IN (subquery)

$\not\subseteq$

NOT EXISTS (subquery)

EXISTS (subquery)

ALL/SOME operator (subquery)

Comparison of sets to each other:

$X \subset Y$

important operations, but not supported directly \otimes

Must rewrite into complicated form

$X - Y = \emptyset \Leftrightarrow X \subseteq Y$

$X = Y$

similar issues, even messier

Set membership



Tuesday, March 10, 2020 9:41 PM

predicate (usually in WHERE clause):

<attribute-list> IN (SELECT <attribute-list> FROM... WHERE ...)

IMPORTANT: attribute lists must be same (compatible)

Saw

userID	title	Release_yr	stars
12345	Little Women	2019	5
12345	Little Women	1994	4
54321	Little Women	2019	4
54321	Finding Dory	2016	3
67890	Little Women	2019	5
67890	Finding Dory	2016	3

A Find people who saw Little Women(2019) and Finding Dory (2016)

SELECT userID FROM Saw

WHERE title = 'LW' AND releaseYR = 2019

AND

userID IN (

SELECT id FROM Saw
WHERE title = 'FD' AND releaseYR = 2016)

predicate

Find people who saw Little Women (2019) and did not see Finding Dory (2016)

IN → NOT IN

Set difference

Aside: intersections with joins:

Tuesday, March 10, 2020 9:50 PM

Saw 1			
userID	title	Release_yr	stars
12345	Little Women	2019	5
12345	Little Women	1994	4
54321	Little Women	2019	4
54321	Finding Dory	2016	3
67890	Little Women	2019	5
67890	Finding Dory	2016	3

Saw 2			
userID	title	Release_yr	stars
12345	Little Women	2019	5
12345	Little Women	1994	4
54321	Little Women	2019	4
54321	Finding Dory	2016	3
67890	Little Women	2019	5
67890	Finding Dory	2016	3

Find people who saw Little Women(2019) and Finding Dory (2016)

```

SELECT
FROM
    Saw AS Saw1 JOIN
    Saw AS Saw2
USING (userID)
WHERE
    Saw1.title = 'Lw' AND
    Saw1.Release-yr = 2019
    AND
    Saw2.title = 'FD'
    AND
    Saw2.Release-yr = 2016
  
```

Scalar subqueries

Tuesday, March 10, 2020 9:53 PM

- Query that is guaranteed to return a single row (no matter what valid data is in the database)
- Can be used like a value of the returned type, e.g.
 - Comparison with <, >, =, etc
 - Arithmetic operations if it's a number
- Can be used in comparisons to result of aggregation
 - Find the movie(s) that got the highest rating (from someone)

userID	title	Release_yr	stars
12345	Little Women	2019	5
12345	Little Women	1994	4
54321	Little Women	2019	4
54321	Finding Dory	2016	3
67890	Little Women	2019	5
67890	Finding Dory	2016	3

~~SELECT title, Release_yr, MAX(stars)~~

DISTINCT

Remember that
attributes in
SELECT must
be aggregated
or in GROUP BY

① Find max ~~~~~ Scalar subquery
② " movies that got that score

~~SELECT title, rel_yr~~ DISTINCT

FROM Saw

WHERE numStars =

(SELECT max(numStars)
FROM Saw)

```
(SELECT max (num stars)  
FROM Saw
```

VIEWS

Tuesday, March 10, 2020 9:58 PM

CREATE VIEW view_name AS (query)

Find movie with the highest average score:

1. Find average score of each movie
2. Find max of those
3. Find title/year of the movie whose average is that max

CREATE VIEW avgStars AS

(SELECT title, relYr, AVG(numStars)
AS avs
FROM Saw
GROUP BY title, relYR)

SELECT title, relYR ③
FROM avgStars

WHERE avs = (SELECT max(avs)
FROM avgStars)

Avg Stars		
title	relYr	AVS
LW	2019	4.5
LW	~	3.5
FD	2016	4.5

title relYr

FD 2016

LW 2019

avS

~

4.5

4.5

Max

END OF Recap

Monday, February 13, 2023 4:26 PM



Correlation Variables

- Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
  from section as S
 where semester = 'Fall' and year= 2009 and
       exists (select *
                  from section as T
                 where semester = 'Spring' and year= 2010
                   and S.course_id= T.course_id);
```

- Correlated subquery
- Correlation name or correlation variable
- Note: correlation can severely impact efficiency

- Subquery may use value(s) of attributes from the outer query
 - Useful for looking for pairs of tuples, t1, t2 where there is some relationship between the value of attribute in t1 and in t2
- Find people who gave a higher rating to movie X than to movie Y
 LW 2019 FD 2016
- These can sometimes be done with JOIN

userID	title	Release_yr	stars	num
12345	Little Women	2019	5	
12345	Little Women	1994	4	
54321	Little Women	2019	4	
✓ 54321	Finding Dory	2016	3	54321
67890	Little Women	2019	5	
67890	Finding Dory	2016	3	

SELECT userID FROM Saw AS Saw1
 WHERE title = 'LW' AND Release_yr = 2019
 AND userID IN
 (SELECT userID FROM Saw AS Saw2
 WHERE title = 'FD' AND Release_yr = 2016)

AND

)

Alternative with JOIN :

SELECT — FROM Saw AS Saw1
 JOIN Saw AS Saw2 USING (userID)
 WHERE Saw1.title — — —
 AND Saw2.title — — .

AND Saw1.numStars > Saw2.numStars



Notes on Execution of SQL Queries

- A SQL query produces a table as output
- In fact, every of the steps on previous slide starts with one or more table, and produces a table
- Thus, the meaning of nesting in the **from** clause is obvious
- For nesting in the **where**, **having**, and **select** clauses, keep in mind that these clauses are applied to each row of the table entering it
- So if there are correlation variables, the subqueries may have to be repeatedly executed
- Exception for **group by** in many DBMS: if you group by a primary key, other attributes in that table can be referred to after group-by:

```
select D.dept_name, D.budget, avg(salary)
  from instructor I join department D
 where D.budget > 100000
   group by D.dept_name
```

Set Theory operations

Tuesday, March 10, 2020 9:38 PM

Set theory operation

- Union
- Intersection
- Set Difference
- Belongs to
- Does not belong to
- Is empty $\Rightarrow \emptyset$
- Is not empty $\neq \emptyset$
- Comparison to each element

SQL Predicate

- UNION [can often be done with pred1 OR pred2 in the WHERE clause, instead]
- INTERSECT [Not supported in MySQL]
- EXCEPT [MINUS in Oracle; not supported in MySQL]
- IN (subquery describing X)
- NOT IN (subquery)
- NOT EXISTS (subquery)
- EXISTS (subquery)
- ALL/SOME operator (subquery)

NOT EXISTS = \emptyset
EXISTS $\neq \emptyset$

Comparison of sets to each other:

$X \subset Y$

important operations, but not supported directly \otimes

Must rewrite into complicated form

$$X - Y = \emptyset \Leftrightarrow X \subseteq Y$$

$X = Y$

similar issues, even messier



Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists** $r \Leftrightarrow r \neq \emptyset$
- **not exists** $r \Leftrightarrow r = \emptyset$

- NOT EXISTS (subquery) \Leftrightarrow result of subquery is empty
- This is a predicate
 - True: subquery result is empty
 - False: subquery result is not empty
- EXISTS (subquery) \Leftrightarrow result of subquery is not empty

Note: "... WHERE EXISTS (subquery) can often be done with a join instead. Think again if you're using this construct.



Not Exists

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name
from student as S
where not exists ( (select course_id
                     from course
                     where dept_name = 'Biology')
except
( select T.course_id
  from takes as T
  where S.ID = T.ID));
```

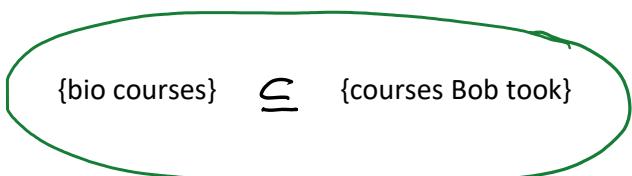
- Note that $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
- Note: Cannot write this query using = all and its variants

Subsets

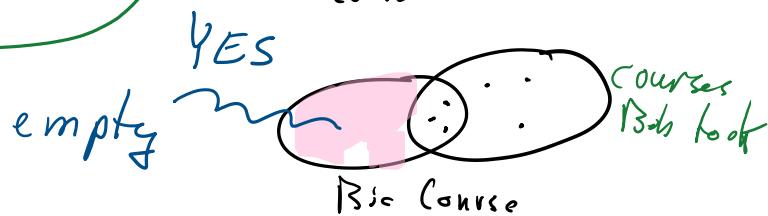
Tuesday, March 10, 2020 10:09 PM

Often want to know whether results of one subquery are all members of the result of another subquery:

Did Bob take **every** Bio course (`course_id LIKE 'Bio%'`) ?



No



YES

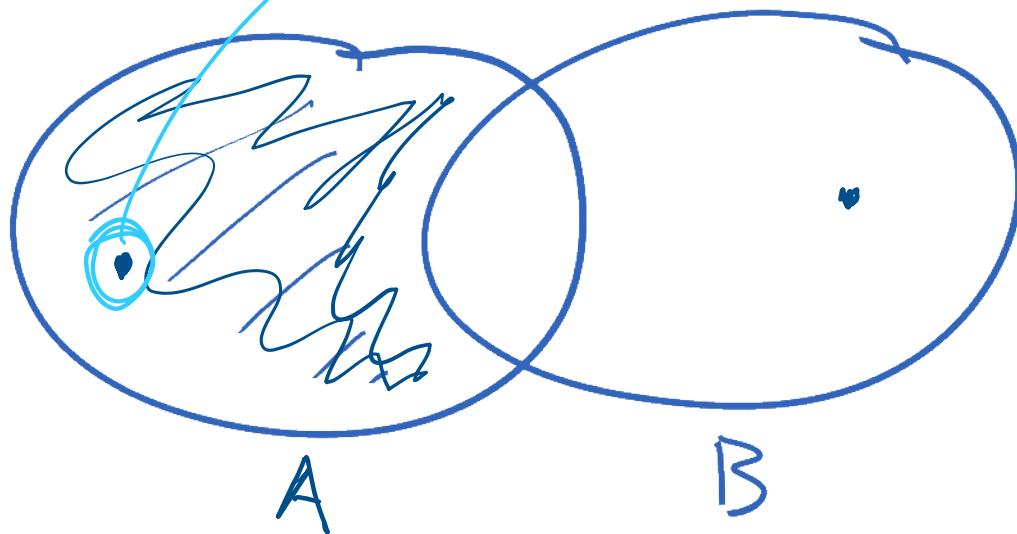
Is $\{ \text{bio courses} \} - \{ \text{courses Bob took} \}$ empty ?

Characterizations of A subset of B:

Tuesday, March 10, 2020

10:12 PM

counter-example



$$A - B = \emptyset \text{ iff } A \subseteq B$$

Which widgets come in every color?

Tuesday, March 10, 2020 10:13 PM

Widgets



Do squares come in
every color **NO!**

Always
empty



Colors of
Square
widgets

All colors

Square widgets
don't come in green

How do this in SQL?

Is $\{\text{all colors}\} - \{\text{colors of squares}\}$
= Ø

NOT

EXISTS

(SELECT colors FROM widgets -- all colors

EXCEPT

(SELECT colors FROM widgets
WHERE shape = 'square') --
of
square
widgets

FALSE

.....

What about triangles



NOT EXISTS

(SELECT colors FROM Widget)

EXCEPT (SELECT colors FROM

widget WHERE

shape = 'triangle'))

General question: Find shapes that come in correlated query that looks at each shape in the outer query every color

SELECT DISTINCT shape FROM widget AS w

WHERE

NOT EXISTS (

(SELECT color FROM widget)

EXCEPT

(SELECT color FROM ...))

(SELECT color FROM widget
WHERE shape = w.shape))

Another approach

Tuesday, October 4, 2022 7:32 PM

$$|\{ \text{bio courses} \}| = |\{ \text{bio courses taken by current student} \}|$$

SELECT s.ID, c.name FROM student AS s

WHERE

(SELECT COUNT(*) FROM course WHERE dept-name = 'Biology')

= (SELECT COUNT(DISTINCT course_id)
FROM takes | NATURAL JOIN course
WHERE s.ID = t.id -- current student
AND dept-name = 'Biology')

$$A \subseteq B$$

Subsets in TRC

Friday, February 23, 2024 11:17 AM

$$x \in A \Rightarrow x \in B$$

$$\{ \text{students } s \mid \begin{array}{l} x \text{ is a bio course} \\ \Rightarrow s \text{ took } x \end{array} \}$$

$$\{ \text{widgets } w \mid \begin{array}{l} c \text{ is a color} \\ \Rightarrow w \text{ comes in color } c \end{array} \}$$

$$\{ t \mid \forall x ((x \in \text{Courses} \wedge x[\text{dept-name}] = 'Bio'))$$

$$\Rightarrow \exists y \in \text{Takes} (y[\text{id}] = t[\text{id}]$$

$$\wedge y[\text{course-id}] = x[\text{course-id}]) \}$$



Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- Find all courses that were offered at most once in 2009

```
select T.course_id
  from course as T
 where unique (select R.course_id
                  from section as R
                 where T.course_id= R.course_id
                   and R.year = 2009);
```



Derived Relations

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
  from (select dept_name, avg(salary) as avg_salary
          from instructor
         group by dept_name)
       where avg_salary > 42000;
```

- Note that we do not need to use the **having** clause
- Another way to write above query

```
select dept_name, avg_salary
  from (select dept_name, avg(salary)
          from instructor
         group by dept_name) as dept_avg(dept_name, avg_salary)
       where avg_salary > 42000;
```



With Clause

- The **with** clause provides a way of defining a temporary view whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```
with max_budget (value) as  
    (select max(budget)  
     from department)  
select budget  
from department, max_budget  
where department.budget = max_budget.value;
```

- Note: often easy to output max, but hard to output who has the max

*definition of
max-budget
dropped*



Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total(dept_name, value) as
  (select dept_name, sum(salary)
   from instructor
   group by dept_name),
dept_total_avg(value) as
  (select avg(value)
   from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value >= dept_total_avg.value;
```



Scalar Subquery

```
select dept_name,  
       (select count(*)  
        from instructor  
       where department.dept_name = instructor.dept_name)  
      as num_instructors  
  from department;
```



Modification of the Database – Deletion

- Delete all instructors

```
delete from instructor
```

- Delete all instructors from the Finance department

```
delete from instructor  
where dept_name= 'Finance';
```

- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

```
delete from instructor  
where dept_name in (select dept_name  
         from department  
         where building = 'Watson');
```



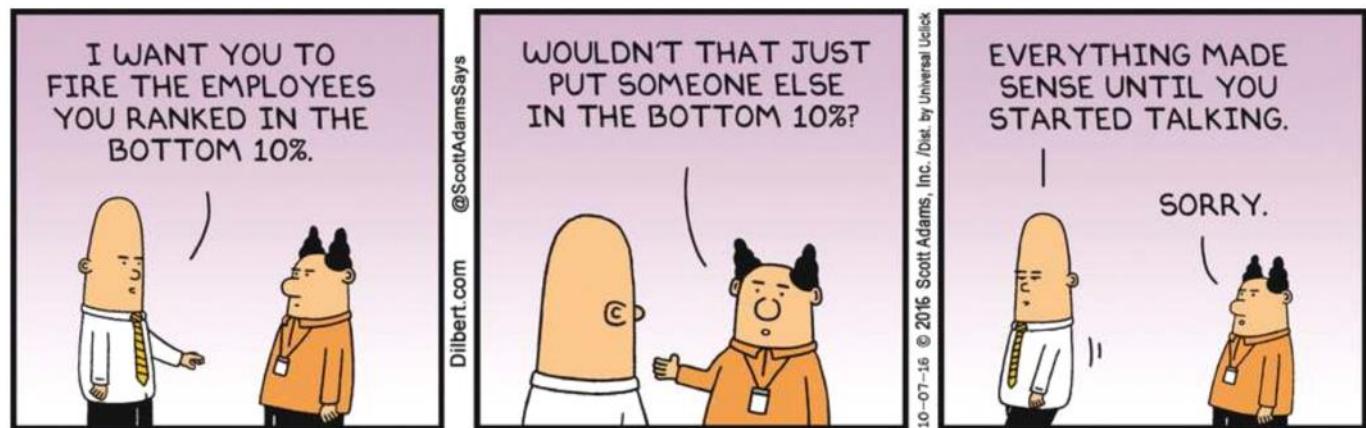
Example Query

- Delete all instructors whose salary is less than the average salary of instructors

delete from *instructor*

where *salary* < (**select avg** (*salary*) **from** *instructor*);

- Problem: as we delete tuples from deposit, the average salary changes --> what would happen?
- Solution used in SQL:
 1. First, compute **avg** salary and find all tuples to delete
 2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)





Modification of the Database – Insertion

- Add a new tuple to *course*

```
insert into course  
    values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

Can insert multiple rows

- or equivalently

```
insert into course (course_id, title, dept_name, credits)  
    values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

(- - - -),
(- - - -)

- Add a new tuple to *student* with *tot_creds* set to null

```
insert into student  
    values ('3003', 'Green', 'Finance', null);
```



Modification of the Database – Insertion

- Add all instructors to the *student* relation with tot_creds set to 0

```
insert into student
    select ID, name, dept_name, 0
        from instructor
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation (otherwise queries like

```
insert into table1 select * from table1
```

would cause problems)



Modification of the Database – Updates

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others receive a 5% raise
 - Write two **update** statements:

```
update instructor  
  set salary = salary * 1.03  
  where salary > 100000;
```

```
update instructor  
  set salary = salary * 1.05  
  where salary <= 100000;
```

- The order is important (why?)
- Can be done better using the **case** statement (next slide)



Case Statement for Conditional Updates

- Same query as before but with case statement

```
update instructor  
set salary = case  
    when salary <= 100000 then salary * 1.05  
    else salary * 1.03  
end
```

Handwritten annotations: 'condition' points to the 'when' clause, and 'action' points to the 'then' or 'else' clause.



Updates with Scalar Subqueries

- Recompute and update tot_creds value for all students

```
update student S
  set tot_cred = (
    select sum(credits)
      from takes natural join course
     where S.ID= takes.ID and
           takes.grade <> 'F' and
           takes.grade is not null);
```

- Sets tot_creds to null for students who have not taken any course
- Instead of **sum(credits)**, use:

```
case
  when sum(credits) is not null then sum(credits)
  else 0
end
```

Comparison of scalar to non-scalar (Boolean aggregation operators)

Tuesday, March 10, 2020 10:00 PM

- ALL (subquery)
- SOME (subquery)
- Can do these other ways, but this is sometimes convenient
- Find instructors whose salaries are higher than *every* Comp Sci instructor

Select ID FROM instructor

WHERE salary >

ALL (SELECT salary FROM instructor
WHERE deptname = 'Comp. Sci.')



Set Comparison

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept name =  
'Biology';
```

- Same query using **> some** clause

```
select name  
from instructor  
where salary > some (select salary  
from instructor  
where dept name = 'Biology');
```

- Note: could also use min here (greater than the minimum)



Definition of Some Clause

- $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F <\text{comp}> t)$
Where $<\text{comp}>$ can be: $<$, \leq , $>$, $=$, \neq

(5 < some	0
	5
	6

) = true

(read: 5 < some tuple in the relation)

(5 < some	0
	5

) = false

(5 = some	0
	5

) = true

(5 \neq some	0
	5

) = true (since $0 \neq 5$)

$(= \text{some}) \equiv \text{in}$

However, $(\neq \text{some}) \not\equiv \text{not in}$



Example Query

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name  
from instructor  
where salary > all (select salary  
         from instructor  
         where dept name = 'Biology');
```



Definition of all Clause

n $F <\text{comp}> \text{all } r \Leftrightarrow \forall t \in r (F <\text{comp}> t)$

$(5 < \text{all} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$

$(5 < \text{all} \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$

$(5 = \text{all} \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 \neq \text{all} \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true} (\text{since } 5 \neq 4 \text{ and } 5 \neq 6)$

$(\neq \text{all}) \equiv \text{not in}$
However, $(= \text{all}) \not\equiv \text{in}$

Lecture 5, continued

Tuesday, October 18, 2022 3:34 PM

- |
 - Outer Joins
 - Views
 - Integrity Constraints
 - Quick Overview of some other topics



Chapter 4: Intermediate SQL

- Join Expressions
- Views
- Transactions
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization



Joined Relations

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join.
- The join operations are typically used as subquery expressions in the **from** clause.



Join operations – Example

- Relation *course*

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation *prereq*

course NATURAL JOIN prereq

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Note: prereq information missing for CS-315 and course information missing for CS-437.

347



Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.

Left Outer Join

- course **natural left outer join prereq**

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

Note: read `prere_id` as `prereq_id`



Right Outer Join

- course **natural right outer join** *prereq*

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null

CS-347 null null null CS-101

Full Outer Join

- course **natural full outer join** *prereq*

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101



Joined Relations

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>
inner join
left outer join
right outer join
full outer join

<i>Join Conditions</i>
natural
on <predicate>
using (A_1, A_1, \dots, A_n)



Joined Relations – Examples

- course **inner join** prereq **on**
 $course.course_id = prereq.course_id$

course_id	title	dept_name	credits	prere_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- course **left outer join** prereq **on**
 $course.course_id = prereq.course_id$

course_id	title	dept_name	credits	prere_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	null	null



Joined Relations – Examples

- course **natural right outer join** prereq

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null

- course **full outer join** prereq **using** (course_id)

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101



Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.



View Definition

- A view is defined using the **create view** statement which has the form

create view *v* as < query expression >

where <query expression> is any legal SQL expression.
The view name is represented by *v*.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
 - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view.



Example Views

- A view of instructors without their salary
create view faculty as
select ID, name, dept_name
from instructor
- Find all instructors in the Biology department
select name
from faculty
where dept_name = 'Biology'
- Create a view of department salary totals
create view departments_total_salary(dept_name, total_salary) as
select dept_name, sum (salary)
from instructor
group by dept_name;



Views Defined Using Other Views

- **create view physics_fall_2009 as**
select course.course_id, sec_id, building, room_number
from course, section
where course.course_id = section.course_id
and course.dept_name = 'Physics'
and section.semester = 'Fall'
and section.year = '2009';

- **create view physics_fall_2009_watson as**
select course_id, room_number
from physics_fall_2009
where building= 'Watson';



Views Defined Using Other Views

- One view may be used in the expression defining another view,
- A view relation v_1 is said to *depend directly* on a view relation v_2 if v_2 is used in the expression defining v_1
- A view relation v_1 is said to *depend on* view relation v_2 if either v_1 depends directly to v_2 or there is a path of dependencies from v_1 to v_2
- A view relation v is said to be *recursive* if it depends on itself.



View Expansion

- A way to define the meaning of views defined in terms of other views.
- Let view v_1 be defined by an expression e_1 that may itself contain uses of view relations.
- View expansion of an expression repeats the following replacement step:
 - repeat**
 - Find any view relation v_i in e_1
 - Replace the view relation v_i by the expression defining v_i
 - until** no more view relations are present in e_1
- As long as the view definitions are not recursive, this loop will terminate.



How Can your DBMS Maintain Views?

- If a view is often used in queries, it might be very inefficient to recompute the view from scratch every time it is used.



How Can your DBMS Maintain Views?

- If a view is often used in queries, it might be very inefficient to recompute the view from scratch every time it is used.
- To deal with this, your DBMS will often materialize and then maintain the view under updates to the underlying tables (“materialized views”)

- Example 1: the faculty view on slide 12:
 - Every insert, delete, update to instructor can be translated into a suitable insert, delete, or update on the materialized view



How Can your DBMS Maintain Views?

- If a view is often used in queries, it might be very inefficient to recompute the view from scratch every time it is used.
- To deal with this, your DBMS will often materialize and then maintain the view under updates to the underlying tables (“materialized views”)

- Example 1: the faculty view on slide 12:
 - Every insert, delete, update to instructor can be translated into a suitable insert, delete, or update on the materialized view
- Example 2: the departments_total_salary view on slide 12:
 - Every insert, delete, and update to instructor can be applied to the view by increasing, decreasing, or updating department salary totals
- The best strategy depends on number of view accesses versus number of updates to the underlying table



Update of a View

- Add a new tuple to *faculty* view which we defined earlier
 - insert into faculty values ('30765', 'Green', 'Music');**
- This insertion could be represented by the insertion of the tuple
 - ('30765', 'Green', 'Music', null)into the *instructor* relation
- Is this legal?



Update of a View

- Add a new tuple to *faculty* view which we defined earlier
 - insert into faculty values ('30765', 'Green', 'Music');**
- This insertion could be represented by the insertion of the tuple
 - ('30765', 'Green', 'Music', null)into the *instructor* relation
- Is this legal?
- Yes, assuming we may assign a null value to salary
- If salary attribute was NOT NULL, cannot do this



Some Updates cannot be Translated Uniquely

- **create view** *instructor_info* **as**
select *ID, name, building*
from *instructor, department*
where *instructor.dept_name= department.dept_name*;
- **insert into** *instructor info* **values** ('69987', 'White', 'Taylor');
 - ▶ which department, if multiple departments in Taylor?
 - ▶ what if no department is in Taylor?
- Most SQL implementations allow updates only on simple views
 - The **from** clause has only one database relation.
 - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
 - Any attribute not listed in the **select** clause can be set to null
 - The query does not have a **group by** or **having** clause.
- But e.g., Oracle allows some updates on views with several relations



And Some Not at All

- ***create view history_instructors as***
select *
from instructor
where dept_name= 'History' ;
- Insert ('25566', 'Brown', 'Biology', 100000) into
history_instructors



Transactions

- Unit of work
- Atomic transaction
 - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions
- Transactions begin implicitly
 - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
 - Can turn off auto commit for a session (e.g. using API)
 - In SQL:1999, can use: **begin atomic end**



Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
 - A checking account must have a balance greater than \$10,000.00.
 - A salary of a bank employee must be at least \$4.00 an hour.
 - A customer must have a (non-null) phone number.



Constraints on a Single Relation

- **not null**
- **primary key**
- **unique**
- **check (P)**, where P is a predicate



single attribute
tuple of a relation



single row

check (salary > 0 and salary < 500000)



Not Null and Unique Constraints

■ not null

- Declare *name* and *budget* to be **not null**

name varchar(20) not null

budget numeric(12,2) not null

■ unique (A_1, A_2, \dots, A_m)

- The unique specification states that the attributes A_1, A_2, \dots, A_m form a candidate key.
- Candidate keys are permitted to be null (in contrast to primary keys).



The check clause

■ **check (P)**

where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (
    course_id varchar (8),
    sec_id varchar (8),
    semester varchar (6),
    year numeric (4,0),
    building varchar (15),
    room_number varchar (7),
    time_slot_id varchar (4),
    primary key (course_id, sec_id, semester, year),
    check (semester in ('Fall', 'Winter', 'Spring',
    'Summer'))
```



Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - Example: If “Biology” is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for “Biology”.
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S.



Cascading Actions in Referential Integrity

- **create table course (**
- course_id char(5) primary key,**
- title varchar(20),**
- dept_name varchar(20) references department**
-)
- **create table course (**
- ...
dept_name varchar(20),
foreign key (dept_name) references department (dept_name)
- on delete cascade**
- on update cascade,**
- ...
)
- alternative actions to cascade: **set null, set default**

What happens when a violation is caused by changing the referenced table (department)

Database System Concepts - 6th Edition
Modified by T. Suel for CS6083, NYU School of Engineering, Spring 2022 4.33

©Silberschatz, Korth and Sudarshan

Cascade : delete the violating tuples from referencing table (course)

Set NULL : change attribute value to NULL

Set DEFAULT



Complex Check Clauses

- **check** (*time_slot_id* in (select *time_slot_id* from *time_slot*))
 - why not use a foreign key here?
- Every section has at least one instructor teaching the section.
 - how to write this?
- Unfortunately: subquery in check clause not supported by pretty much any database
 - Alternative: triggers (later)
- **create assertion** <assertion-name> **check** <predicate>;
 - Also not supported by anyone



Built-in Data Types in SQL

- **date:** Dates, containing a (4 digit) year, month and date
 - Example: **date** ‘2005-7-27’
- **time:** Time of day, in hours, minutes and seconds.
 - Example: **time** ‘09:00:30’ **time** ‘09:00:30.75’
- **timestamp:** date plus time of day
 - Example: **timestamp** ‘2005-7-27 09:00:30.75’
- **interval:** period of time
 - Example: **interval** ‘1’ day
 - Subtracting a date/time/timestamp value from another gives an interval value
 - Interval values can be added to date/time/timestamp values



Other Features

- **create table** *student*
*(ID varchar (5),
name varchar (20) not null,
dept_name varchar (20),
tot_cred numeric (3,0) default 0,
primary key (ID))*
- **create index** *studentID* index on *student(ID)*
- Large objects
 - *book review clob(10KB)*
 - *image blob(10MB), movie blob(2GB)*
- **create type** construct in SQL creates user-defined type
create type *Dollars as numeric (12,2) final*
 - **create table** *department*
*(dept_name varchar (20),
building varchar (15),
budget Dollars);*



- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person_name char(20) not null
```

- Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.
- **create domain degree_level varchar(10)**
constraint degree_level_test
check (value in ('Bachelors', 'Masters', 'Doctorate'));
- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
 - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
 - **clob**: character large object -- object is a large collection of char data
 - When a query returns a large object, a pointer is returned rather than the large object itself.